



Cloud

Embedding

Vision

Tools

Thinking

Popular



kimi-k2.5

Kimi K2.5 is an open-source, native multimodal agentic model that seamlessly integrates vision and language understanding with advanced agentic capabilities, instant and thinking modes, as well as conversational and agentic paradigms.

cloud

↓ 6,815 Pulls ↗ 1 Tag ⌚ Updated yesterday

glm-4.7-flash

As the strongest model in the 30B class, GLM-4.7-Flash offers a new option for lightweight deployment that balances performance and efficiency.

tools

thinking

↓ 71.1K Pulls ↗ 4 Tags ⌚ Updated 4 days ago

translategemma

A new collection of open translation models built on Gemma 3, helping people communicate across 55 languages.

vision

4b

12b

27b

↓ 90K Pulls ↗ 13 Tags ⌚ Updated 1 week ago

lfm2.5-thinking

LFM2.5 is a new family of hybrid models designed for on-device deployment.

1.2b

↓ 13.8K Pulls ↗ 5 Tags ⌚ Updated 1 week ago

qwen3-vl

The most powerful vision-language model in the Qwen model family to date.

vision

tools

thinking

cloud

2b

4b

8b

30b

32b

235b

↓ 1.2M Pulls ↗ 59 Tags ⌚ Updated 3 months ago

rnj-1

Rnj-1 is a family of 8B parameter open-weight, dense models trained from scratch by Essential AI, optimized for code and STEM with capabilities on par with SOTA open-weight models.

tools

cloud

8b

↓ 299.3K Pulls ↗ 6 Tags ⌚ Updated 1 month ago

ministral-3

The Ministral 3 family is designed for edge deployment, capable of running on a wide range of hardware.

vision

tools

cloud

3b

8b

14b

↓ 324.8K Pulls ↗ 16 Tags ⌚ Updated 1 month ago

qwen3-next

The first installment in the Qwen3-Next series with strong performance in terms of both parameter efficiency and inference speed.

tools

thinking

cloud

80b

↓ 281.8K Pulls ↗ 10 Tags ⌚ Updated 1 month ago

granite4

Granite 4 features improved instruction following (IF) and tool-calling capabilities, making them more effective in enterprise applications.

tools

350m

1b

3b

↓ 594.3K Pulls ↗ 17 Tags ⌚ Updated 3 months ago

nemotron-3-nano

Nemotron 3 Nano - A new Standard for Efficient, Open, and Intelligent Agentic Models

tools

thinking

cloud

30b

↓ 135.3K Pulls ↗ 6 Tags ⌚ Updated 1 month ago

qwen3-embedding

Building upon the foundational models of the Qwen3 series, Qwen3 Embedding provides a comprehensive range of text embeddings models in various sizes

embedding

0.6b

4b

8b

↓ 400.8K Pulls ↗ 12 Tags ⌚ Updated 4 months ago

embeddinggemma

EmbeddingGemma is a 300M parameter embedding model from Google.

embedding

300m

↓ 452K Pulls ↗ 5 Tags ⌚ Updated 4 months ago

devstral-small-2

24B model that excels at using tools to explore codebases, editing multiple files and power software engineering agents.

vision

tools

cloud

24b

↓ 117K Pulls ↗ 6 Tags ⌚ Updated 1 month ago

olmo-3

Olmo is a series of Open language models designed to enable the science of language models. These models are pre-trained on the Dolma 3 dataset and post-trained on the Dolci datasets.

7b

32b

↓ 108.7K Pulls ↗ 15 Tags ⌚ Updated 1 month ago

deepseek-ocr

DeepSeek-OCR is a vision-language model that can perform token-efficient OCR.

vision 3b

↓ 130.6K Pulls ↗ 3 Tags ⌚ Updated 2 months ago

deepseek-v3.1

DeepSeek-V3.1-Terminus is a hybrid model that supports both thinking mode and non-thinking mode.

tools thinking cloud 671b

↓ 298.8K Pulls ↗ 8 Tags ⌚ Updated 4 months ago

olmo-3.1

Olmo is a series of Open language models designed to enable the science of language models. These models are pre-trained on the Dolma 3 dataset and post-trained on the Dolci datasets.

tools 32b

↓ 64.5K Pulls ↗ 10 Tags ⌚ Updated 1 month ago

gemini-3-pro-preview

Google's most intelligent model with SOTA reasoning and multimodal understanding, and powerful agentic and vibe coding capabilities.

cloud

↓ 92.4K Pulls ↗ 1 Tag ⌚ Updated 2 months ago

functiongemma

FunctionGemma is a specialized version of Google's Gemma 3 270M model fine-tuned explicitly for function calling.

tools 270m

↓ 51K Pulls ↗ 4 Tags ⌚ Updated 1 month ago

devstral-2

123B model that excels at using tools to explore codebases, editing multiple files and power software engineering agents.

tools

cloud

123b

↓ 59.7K Pulls ↗ 6 Tags ⌚ Updated 1 month ago

gemini-3-flash-preview

Gemini 3 Flash offers frontier intelligence built for speed at a fraction of the cost.

cloud

↓ 46.8K Pulls ↗ 2 Tags ⌚ Updated 1 month ago

cogito-2.1

The Cogito v2.1 LLMs are instruction tuned generative models. All models are released under MIT license for commercial use.

cloud

671b

↓ 53.5K Pulls ↗ 6 Tags ⌚ Updated 2 months ago

glm-4.7

Advancing the Coding Capability

cloud

↓ 23.7K Pulls ↗ 1 Tag ⌚ Updated 1 month ago

gpt-oss-safeguard

gpt-oss-safeguard-20b and gpt-oss-safeguard-120b are safety reasoning models built-upon gpt-oss

tools

thinking

20b

120b

↓ 55.1K Pulls ↗ 3 Tags ⌚ Updated 3 months ago

nomic-embed-text-v2-moe

nomic-embed-text-v2-moe is a multilingual MoE text embedding model that excels at multilingual retrieval.

embedding

↓ 25.5K Pulls 1 Tag Updated 1 month ago

glm-4.6

Advanced agentic, reasoning and coding capabilities.

cloud

↓ 53.3K Pulls 1 Tag Updated 3 months ago

minimax-m2

MiniMax M2 is a high-efficiency large language model built for coding and agentic workflows.

cloud

↓ 44.9K Pulls 1 Tag Updated 3 months ago

deepseek-v3.2

DeepSeek-V3.2, a model that harmonizes high computational efficiency with superior reasoning and agent performance.

cloud

↓ 18.6K Pulls 1 Tag Updated 1 month ago

minimax-m2.1

Exceptional multilingual capabilities to elevate code engineering

cloud

↓ 11.9K Pulls 1 Tag Updated 1 month ago

kimi-k2-thinking

Kimi K2 Thinking, Moonshot AI's best open-source thinking model.

cloud

↓ 24.5K Pulls 1 Tag Updated 2 months ago

kimi-k2

A state-of-the-art mixture-of-experts (MoE) language model. Kimi K2-Instruct-0905 demonstrates significant improvements in performance on public benchmarks and real-world coding agent tasks.

cloud

↓ 33.5K Pulls 1 Tag Updated 4 months ago

mistral-large-3

A general-purpose multimodal mixture-of-experts model for production-grade tasks and enterprise workloads.

cloud

↓ 13.6K Pulls 1 Tag Updated 1 month ago

gemma3

The current, most capable model that runs on a single GPU.

vision

cloud

270m

1b

4b

12b

27b

↓ 30.9M Pulls 29 Tags Updated 1 month ago

qwen3

Qwen3 is the latest generation of large language models in Qwen series, offering a comprehensive suite of dense and mixture-of-experts (MoE) models.

tools

thinking

0.6b

1.7b

4b

8b

14b

30b

32b

235b

↓ 18.2M Pulls 58 Tags Updated 3 months ago

gpt-oss

OpenAI's open-weight models designed for powerful reasoning, agentic tasks, and versatile developer use cases.

tools thinking cloud 20b 120b

↓ 6.3M Pulls ↗ 5 Tags ⌚ Updated 3 months ago

qwen3-coder

Alibaba's performant long context models for agentic and coding tasks.

tools cloud 30b 480b

↓ 2.7M Pulls ↗ 10 Tags ⌚ Updated 4 months ago

mistral-small3.2

An update to Mistral Small that improves on function calling, instruction following, and less repetition errors.

vision tools 24b

↓ 1.2M Pulls ↗ 5 Tags ⌚ Updated 7 months ago

gemma3n

Gemma 3n models are designed for efficient execution on everyday devices such as laptops, tablets or phones.

e2b e4b

↓ 1.1M Pulls ↗ 9 Tags ⌚ Updated 7 months ago

qwen2.5vl

Flagship vision-language model of Qwen and also a significant leap from the previous Qwen2-VL.

vision 3b 7b 32b 72b

↓ 1.2M Pulls ↗ 17 Tags ⌚ Updated 8 months ago

magistral

Magistral is a small, efficient reasoning model with 24B parameters.

tools thinking 24b

↓ 1M Pulls ↗ 5 Tags ⌚ Updated 7 months ago

cogito

Cogito v1 Preview is a family of hybrid reasoning models by Deep Cogito that outperform the best available open models of the same size, including counterparts from LLaMA, DeepSeek, and Qwen across most standard benchmarks.

tools 3b 8b 14b 32b 70b

↓ 1.3M Pulls ↗ 20 Tags ⌚ Updated 9 months ago

phi4-reasoning

Phi 4 reasoning and reasoning plus are 14-billion parameter open-weight reasoning models that rival much larger models on complex reasoning tasks.

14b

↓ 1.1M Pulls ↗ 9 Tags ⌚ Updated 9 months ago

llama4

Meta's latest collection of multimodal models.

vision tools 16x17b 128x17b

↓ 1.1M Pulls ↗ 11 Tags ⌚ Updated 7 months ago

granite3.3

IBM Granite 2B and 8B models are 128K context length language models that have been fine-tuned for improved reasoning and instruction-following capabilities.

tools 2b 8b

↓ 853.9K Pulls ↗ 3 Tags ⌚ Updated 9 months ago

deepscaler

A fine-tuned version of Deepseek-R1-Distilled-Qwen-1.5B that surpasses the performance of OpenAI's o1-preview with just 1.5B parameters on popular math evaluations.

1.5b

↓ 1.1M Pulls 5 Tags Updated 11 months ago

devstral

Devstral: the best open source model for coding agents

tools

24b

↓ 681.3K Pulls 5 Tags Updated 6 months ago

phi4-mini

Phi-4-mini brings significant enhancements in multilingual support, reasoning, and mathematics, and now, the long-awaited function calling feature is finally supported.

tools

3.8b

↓ 803.9K Pulls 5 Tags Updated 11 months ago

granite3.2-vision

A compact and efficient vision-language model, specifically designed for visual document understanding, enabling automated content extraction from tables, charts, infographics, plots, diagrams, and more.

vision

tools

2b

↓ 689.8K Pulls 5 Tags Updated 11 months ago

deepcoder

DeepCoder is a fully open-Source 14B coder model at O3-mini level, with a 1.5B version also available.

1.5b

14b

↓ 577.8K Pulls 9 Tags Updated 9 months ago

openthinker

A fully open-source family of reasoning models built using a dataset derived by distilling DeepSeek-R1.

7b

32b

↓ 736.6K Pulls ↗ 15 Tags ⌚ Updated 9 months ago

mistral-small3.1

Building upon Mistral Small 3, Mistral Small 3.1 (2503) adds state-of-the-art vision understanding and enhances long context capabilities up to 128k tokens without compromising text performance.

vision

tools

24b

↓ 557K Pulls ↗ 5 Tags ⌚ Updated 9 months ago

exaone-deep

EXAONE Deep exhibits superior capabilities in various reasoning tasks including math and coding benchmarks, ranging from 2.4B to 32B parameters developed and released by LG AI Research.

2.4b

7.8b

32b

↓ 427.5K Pulls ↗ 13 Tags ⌚ Updated 10 months ago

granite3.2

Granite-3.2 is a family of long-context AI models from IBM Granite fine-tuned for thinking capabilities.

tools

2b

8b

↓ 241.1K Pulls ↗ 9 Tags ⌚ Updated 11 months ago

r1-1776

A version of the DeepSeek-R1 model that has been post trained to provide unbiased, accurate, and factual information by Perplexity.

70b

671b

↓ 214K Pulls ↗ 9 Tags ⌚ Updated 11 months ago

phi4-mini-reasoning

Phi 4 mini reasoning is a lightweight open model that balances efficiency with advanced reasoning ability.

3.8b

↓ 132K Pulls ↗ 5 Tags ⌚ Updated 9 months ago

command-a

111 billion parameter model optimized for demanding enterprises that require fast, secure, and high-quality AI

tools

111b

↓ 113.4K Pulls ↗ 5 Tags ⌚ Updated 10 months ago

command-r7b-arabic

A new state-of-the-art version of the lightweight Command R7B model that excels in advanced Arabic language capabilities for enterprises in the Middle East and Northern Africa.

tools

7b

↓ 70.5K Pulls ↗ 5 Tags ⌚ Updated 11 months ago

deepseek-r1

DeepSeek-R1 is a family of open reasoning models with performance approaching that of leading models, such as O3 and Gemini 2.5 Pro.

tools

thinking

1.5b

7b

8b

14b

32b

70b

671b

↓ 77.2M Pulls ↗ 35 Tags ⌚ Updated 7 months ago

llama3.1

Llama 3.1 is a new state-of-the-art model from Meta available in 8B, 70B and 405B parameter sizes.

tools

8b

70b

405b

↓ 109.3M Pulls ↗ 93 Tags ⌚ Updated 1 year ago

llama3.2

Meta's Llama 3.2 goes small with 1B and 3B models.

tools 1b 3b

↓ 54.7M Pulls ↗ 63 Tags ⌚ Updated 1 year ago

nomic-embed-text

A high-performing open embedding model with a large token context window.

embedding

↓ 52.4M Pulls ↗ 3 Tags ⌚ Updated 1 year ago

qwen2.5

Qwen2.5 models are pretrained on Alibaba's latest large-scale dataset, encompassing up to 18 trillion tokens. The model supports up to 128K tokens and has multilingual support.

tools 0.5b 1.5b 3b 7b 14b 32b 72b

↓ 19.8M Pulls ↗ 133 Tags ⌚ Updated 1 year ago

gemma2

Google Gemma 2 is a high-performing and efficient model available in three sizes: 2B, 9B, and 27B.

2b 9b 27b

↓ 14.4M Pulls ↗ 94 Tags ⌚ Updated 1 year ago

qwen2.5-coder

The latest series of Code-Specific Qwen models, with significant improvements in code generation, code reasoning, and code fixing.

tools 0.5b 1.5b 3b 7b 14b 32b

↓ 10.6M Pulls ↗ 199 Tags ⌚ Updated 8 months ago

phi4

Phi-4 is a 14B parameter, state-of-the-art open model from Microsoft.

14b

↓ 7M Pulls 🔖 5 Tags ⌚ Updated 1 year ago

phi3

Phi-3 is a family of lightweight 3B (Mini) and 14B (Medium) state-of-the-art open models by Microsoft.

3.8b

14b

↓ 15.8M Pulls 🔖 72 Tags ⌚ Updated 1 year ago

mistral

The 7B model released by Mistral AI, updated to version 0.3.

tools

7b

↓ 24.6M Pulls 🔖 84 Tags ⌚ Updated 6 months ago

llama3

Meta Llama 3: The most capable openly available LLM to date

8b

70b

↓ 14.4M Pulls 🔖 68 Tags ⌚ Updated 1 year ago

llava

🏠 LLaVA is a novel end-to-end trained large multimodal model that combines a vision encoder and Vicuna for general-purpose visual and language understanding. Updated to version 1.6.

vision

7b

13b

34b

↓ 12.7M Pulls 🔖 98 Tags ⌚ Updated 1 year ago

dolphin3

Dolphin 3.0 Llama 3.1 8B 🐬 is the next generation of the Dolphin series of instruct-tuned models designed to be the ultimate general purpose local model, enabling coding, math, agentic, function calling, and general use cases.

8b

↓ 3.6M Pulls 🏷 5 Tags ⌚ Updated 1 year ago

olmo2

OLMo 2 is a new family of 7B and 13B models trained on up to 5T tokens. These models are on par with or better than equivalently sized fully open models, and competitive with open-weight models such as Llama 3.1 on English academic benchmarks.

7b

13b

↓ 3.5M Pulls 🏷 9 Tags ⌚ Updated 1 year ago

deepseek-v3

A strong Mixture-of-Experts (MoE) language model with 671B total parameters with 37B activated for each token.

671b

↓ 3.3M Pulls 🏷 5 Tags ⌚ Updated 1 year ago

minicpm-v

A series of multimodal LLMs (MLLMs) designed for vision-language understanding.

vision

8b

↓ 4.5M Pulls 🏷 17 Tags ⌚ Updated 1 year ago

llama3.2-vision

Llama 3.2 Vision is a collection of instruction-tuned image reasoning generative models in 11B and 90B sizes.

vision

11b

90b

↓ 3.7M Pulls 🏷 9 Tags ⌚ Updated 8 months ago

mxbai-embed-large

State-of-the-art large embedding model from mixedbread.ai

embedding

335m

↓ 6.9M Pulls ↗ 4 Tags ⌚ Updated 1 year ago

llama3.3

New state of the art 70B model. Llama 3.3 70B offers similar performance compared to the Llama 3.1 405B model.

tools

70b

↓ 3.1M Pulls ↗ 14 Tags ⌚ Updated 1 year ago

qwen2

Qwen2 is a new series of large language models from Alibaba group

tools

0.5b

1.5b

7b

72b

↓ 4.8M Pulls ↗ 97 Tags ⌚ Updated 1 year ago

gemma

Gemma is a family of lightweight, state-of-the-art open models built by Google DeepMind. Updated to version 1.1

2b

7b

↓ 5.8M Pulls ↗ 102 Tags ⌚ Updated 1 year ago

qwen

Qwen 1.5 is a series of large language models by Alibaba Cloud spanning from 0.5B to 110B parameters

0.5b

1.8b

4b

7b

14b

32b

72b

110b

↓ 5.4M Pulls ↗ 379 Tags ⌚ Updated 1 year ago

granite3.1-moe

The IBM Granite 1B and 3B models are long-context mixture of experts (MoE) Granite models from IBM designed for low latency usage.

tools 1b 3b

↓ 2.1M Pulls ↗ 33 Tags ⌚ Updated 1 year ago

falcon3

A family of efficient AI models under 10B parameters performant in science, math, and coding through innovative training techniques.

1b 3b 7b 10b

↓ 2.1M Pulls ↗ 17 Tags ⌚ Updated 1 year ago

mistral-nemo

A state-of-the-art 12B model with 128k context length, built by Mistral AI in collaboration with NVIDIA.

tools 12b

↓ 3.3M Pulls ↗ 17 Tags ⌚ Updated 6 months ago

bge-m3

BGE-M3 is a new model from BAAI distinguished for its versatility in Multi-Functionality, Multi-Linguality, and Multi-Granularity.

embedding 567m

↓ 3.1M Pulls ↗ 3 Tags ⌚ Updated 1 year ago

smollm2

SmolLM2 is a family of compact language models available in three size: 135M, 360M, and 1.7B parameters.

tools 135m 360m 1.7b

↓ 2.4M Pulls ↗ 49 Tags ⌚ Updated 1 year ago

qwq

QwQ is the reasoning model of the Qwen series.

[tools](#)[32b](#)

↓ 2M Pulls ↗ 8 Tags ⌚ Updated 10 months ago

mistral-small

Mistral Small 3 sets a new benchmark in the “small” Large Language Models category below 70B.

[tools](#)[22b](#)[24b](#)

↓ 2.3M Pulls ↗ 21 Tags ⌚ Updated 12 months ago

llama2

Llama 2 is a collection of foundation language models ranging from 7B to 70B parameters.

[7b](#)[13b](#)[70b](#)

↓ 5.2M Pulls ↗ 102 Tags ⌚ Updated 2 years ago

tinyllama

The TinyLlama project is an open endeavor to train a compact 1.1B Llama model on 3 trillion tokens.

[1.1b](#)

↓ 3.5M Pulls ↗ 36 Tags ⌚ Updated 2 years ago

codellama

A large language model that can use text prompts to generate and discuss code.

[7b](#)[13b](#)[34b](#)[70b](#)

↓ 4.1M Pulls ↗ 199 Tags ⌚ Updated 1 year ago

llava-llama3

A LLaVA model fine-tuned from Llama 3 Instruct with better scores in several benchmarks.

vision 8b

↓ 2.1M Pulls ↗ 4 Tags ⌚ Updated 1 year ago

deepseek-coder

DeepSeek Coder is a capable coding model trained on two trillion code and natural language tokens.

1.3b 6.7b 33b

↓ 2.8M Pulls ↗ 102 Tags ⌚ Updated 2 years ago

codegemma

CodeGemma is a collection of powerful, lightweight models that can perform a variety of coding tasks like fill-in-the-middle code completion, code generation, natural language understanding, mathematical reasoning, and instruction following.

2b 7b

↓ 2.1M Pulls ↗ 85 Tags ⌚ Updated 1 year ago

all-minilm

Embedding models on very large sentence level datasets.

embedding 22m 33m

↓ 2.3M Pulls ↗ 10 Tags ⌚ Updated 1 year ago

snowflake-arctic-embed

A suite of text embedding models by Snowflake, optimized for performance.

embedding 22m 33m 110m 137m 335m

↓ 1.8M Pulls ↗ 16 Tags ⌚ Updated 1 year ago

starcoder2

StarCoder2 is the next generation of transparently trained open code LLMs that comes in three sizes: 3B, 7B and 15B parameters.

3b 7b 15b

↓ 1.9M Pulls 🏷 67 Tags ⌚ Updated 1 year ago

deepseek-coder-v2

An open-source Mixture-of-Experts code language model that achieves performance comparable to GPT4-Turbo in code-specific tasks.

16b 236b

↓ 1.5M Pulls 🏷 64 Tags ⌚ Updated 1 year ago

mixtral

A set of Mixture of Experts (MoE) model with open weights by Mistral AI in 8x7b and 8x22b parameter sizes.

tools 8x7b 8x22b

↓ 1.7M Pulls 🏷 70 Tags ⌚ Updated 1 year ago

orca-mini

A general-purpose model ranging from 3 billion parameters to 70 billion, suitable for entry-level hardware.

3b 7b 13b 70b

↓ 1.8M Pulls 🏷 119 Tags ⌚ Updated 2 years ago

smollm

🔥 A family of small models with 135M, 360M, and 1.7B parameters, trained on a new high-quality dataset.

135m 360m 1.7b

↓ 826.4K Pulls 🏷 94 Tags ⌚ Updated 1 year ago

llama2-uncensored

Uncensored Llama 2 model by George Sung and Jarrad Hope.

[7b](#)[70b](#)

↓ 1.7M Pulls ↗ 34 Tags ⌚ Updated 2 years ago

dolphin-phi

2.7B uncensored Dolphin model by Eric Hartford, based on the Phi language model by Microsoft Research.

[2.7b](#)

↓ 1.2M Pulls ↗ 15 Tags ⌚ Updated 2 years ago

dolphin-llama3

Dolphin 2.9 is a new model with 8B and 70B sizes by Eric Hartford based on Llama 3 that has a variety of instruction, conversational, and coding skills.

[8b](#)[70b](#)

↓ 894.5K Pulls ↗ 53 Tags ⌚ Updated 1 year ago

codestral

Codestral is Mistral AI's first-ever code model designed for code generation tasks.

[22b](#)

↓ 743K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

phi

Phi-2: a 2.7B language model by Microsoft Research that demonstrates outstanding reasoning and language understanding capabilities.

[2.7b](#)

↓ 955.7K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

dolphin-mixtral

Uncensored, 8x7b and 8x22b fine-tuned models based on the Mixtral mixture of experts models that excels at coding tasks. Created by Eric Hartford.

8x7b 8x22b

↓ 946.2K Pulls ↗ 70 Tags ⌚ Updated 1 year ago

hermes3

Hermes 3 is the latest version of the flagship Hermes series of LLMs by Nous Research

tools 3b 8b 70b 405b

↓ 533.1K Pulls ↗ 65 Tags ⌚ Updated 1 year ago

phi3.5

A lightweight AI model with 3.8 billion parameters with performance overtaking similarly and larger sized models.

3.8b

↓ 487.8K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

granite-code

A family of open foundation models by IBM for Code Intelligence

3b 8b 20b 34b

↓ 592.8K Pulls ↗ 162 Tags ⌚ Updated 1 year ago

opencoder

OpenCoder is an open and reproducible code LLM family which includes 1.5B and 8B models, supporting chat in English and Chinese languages.

1.5b 8b

↓ 365.3K Pulls ↗ 9 Tags ⌚ Updated 1 year ago

wizardlm2

State of the art large language model from Microsoft AI with improved performance on complex chat, multilingual, reasoning and agent use cases.

7b 8x22b

↓ 598.1K Pulls ↗ 22 Tags ⌚ Updated 1 year ago

moondream

moondream2 is a small vision language model designed to run efficiently on edge devices.

vision 1.8b

↓ 572.8K Pulls ↗ 18 Tags ⌚ Updated 1 year ago

command-r

Command R is a Large Language Model optimized for conversational interaction and long context tasks.

tools 35b

↓ 607.5K Pulls ↗ 32 Tags ⌚ Updated 1 year ago

mistral-large

Mistral Large 2 is Mistral's new flagship model that is significantly more capable in code generation, mathematics, and reasoning with 128k context window and support for dozens of languages.

tools 123b

↓ 460.7K Pulls ↗ 32 Tags ⌚ Updated 1 year ago

granite3.1-dense

The IBM Granite 2B and 8B models are text-only dense LLMs trained on over 12 trillion tokens of data, demonstrated significant improvements over their predecessors in performance and speed in IBM's initial testing.

tools 2b 8b

↓ 308.6K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

dolphin-mistral

The uncensored Dolphin model based on Mistral that excels at coding tasks. Updated to version 2.8.

7b

↓ 633.9K Pulls ↗ 120 Tags ⌚ Updated 1 year ago

paraphrase-multilingual

Sentence-transformers model that can be used for tasks like clustering or semantic search.

embedding

278m

↓ 378.2K Pulls ↗ 3 Tags ⌚ Updated 1 year ago

qwen2-math

Qwen2 Math is a series of specialized math language models built upon the Qwen2 LLMs, which significantly outperforms the mathematical capabilities of open-source models and even closed-source models (e.g., GPT4o).

1.5b

7b

72b

↓ 356.6K Pulls ↗ 52 Tags ⌚ Updated 1 year ago

granite3-dense

The IBM Granite 2B and 8B models are designed to support tool-based use cases and support for retrieval augmented generation (RAG), streamlining code generation, translation and bug fixing.

tools

2b

8b

↓ 309.4K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

yi-coder

Yi-Coder is a series of open-source code language models that delivers state-of-the-art coding performance with fewer than 10 billion parameters.

1.5b

9b

↓ 314.4K Pulls ↗ 67 Tags ⌚ Updated 1 year ago

deepseek-v2

A strong, economical, and efficient Mixture-of-Experts language model.

16b 236b

↓ 390.2K Pulls ↗ 34 Tags ⌚ Updated 1 year ago

aya-expense

Cohere For AI's language models trained to perform well across 23 different languages.

tools 8b 32b

↓ 273.5K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

glm4

A strong multi-lingual general language model with competitive performance to Llama 3.

9b

↓ 355K Pulls ↗ 32 Tags ⌚ Updated 1 year ago

llama-guard3

Llama Guard 3 is a series of models fine-tuned for content safety classification of LLM inputs and responses.

1b 8b

↓ 279.8K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

granite3-moe

The IBM Granite 1B and 3B models are the first mixture of experts (MoE) Granite models from IBM designed for low latency usage.

tools 1b 3b

↓ 270.4K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

snowflake-arctic-embed2

Snowflake's frontier embedding model. Arctic Embed 2.0 adds multilingual support without sacrificing English performance or scalability.

[embedding](#)[568m](#)

↓ 234.6K Pulls ↗ 3 Tags ⌚ Updated 1 year ago

exaone3.5

EXAONE 3.5 is a collection of instruction-tuned bilingual (English and Korean) generative models ranging from 2.4B to 32B parameters, developed and released by LG AI Research.

[2.4b](#)[7.8b](#)[32b](#)

↓ 220K Pulls ↗ 13 Tags ⌚ Updated 1 year ago

aya

Aya 23, released by Cohere, is a new family of state-of-the-art, multilingual models that support 23 languages.

[8b](#)[35b](#)

↓ 350.1K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

yi

Yi 1.5 is a high-performing, bilingual language model.

[6b](#)[9b](#)[34b](#)

↓ 492.2K Pulls ↗ 174 Tags ⌚ Updated 1 year ago

bakllava

BakLLaVA is a multimodal model consisting of the Mistral 7B base model augmented with the LLaVA architecture.

[vision](#)[7b](#)

↓ 465.2K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

codeqwen

CodeQwen1.5 is a large language model pretrained on a large amount of code data.

7b

↓ 368.7K Pulls ↗ 30 Tags ⌚ Updated 1 year ago

codegeex4

A versatile model for AI software development scenarios, including code completion.

9b

↓ 302.4K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

shieldgemma

ShieldGemma is set of instruction tuned models for evaluating the safety of text prompt input and text output responses against a set of defined safety policies.

2b

9b

27b

↓ 240.6K Pulls ↗ 49 Tags ⌚ Updated 1 year ago

athene-v2

Athene-V2 is a 72B parameter model which excels at code completion, mathematics, and log extraction tasks.

tools

72b

↓ 214.6K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

llama3-groq-tool-use

A series of models from Groq that represent a significant advancement in open-source AI capabilities for tool use/function calling.

tools

8b

70b

↓ 284.3K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

zephyr

Zephyr is a series of fine-tuned versions of the Mistral and Mixtral models that are trained to act as helpful assistants.

[7b](#)[141b](#)

↓ 474.4K Pulls ↗ 40 Tags ⌚ Updated 1 year ago

reader-lm

A series of models that convert HTML content to Markdown content, which is useful for content conversion tasks.

[0.5b](#)[1.5b](#)

↓ 245.5K Pulls ↗ 33 Tags ⌚ Updated 1 year ago

internlm2

InternLM2.5 is a 7B parameter model tailored for practical scenarios with outstanding reasoning capability.

[1m](#)[1.8b](#)[7b](#)[20b](#)

↓ 288.4K Pulls ↗ 65 Tags ⌚ Updated 1 year ago

granite-embedding

The IBM Granite Embedding 30M and 278M models are text-only dense biencoder embedding models, with 30M available in English only and 278M serving multilingual use cases.

[embedding](#)[30m](#)[278m](#)

↓ 189.3K Pulls ↗ 6 Tags ⌚ Updated 1 year ago

llama3-chatqa

A model from NVIDIA based on Llama 3 that excels at conversational question answering (QA) and retrieval-augmented generation (RAG).

[8b](#)[70b](#)

↓ 312.1K Pulls ↗ 35 Tags ⌚ Updated 1 year ago

nemotron

Llama-3.1-Nemotron-70B-Instruct is a large language model customized by NVIDIA to improve the helpfulness of LLM generated responses to user queries.

[tools](#)[70b](#)

↓ 211.4K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

reflection

A high-performing model trained with a new technique called Reflection-tuning that teaches a LLM to detect mistakes in its reasoning and correct course.

[70b](#)

↓ 233.8K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

nemotron-mini

A commercial-friendly small language model by NVIDIA optimized for roleplay, RAG QA, and function calling.

[tools](#)[4b](#)

↓ 225.9K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

command-r-plus

Command R+ is a powerful, scalable large language model purpose-built to excel at real-world enterprise use cases.

[tools](#)[104b](#)

↓ 316.9K Pulls ↗ 21 Tags ⌚ Updated 1 year ago

deepseek-llm

An advanced language model crafted with 2 trillion bilingual tokens.

[7b](#)[67b](#)

↓ 396.8K Pulls ↗ 64 Tags ⌚ Updated 2 years ago

llama3-gradient

This model extends LLama-3 8B's context length from 8k to over 1m tokens.

[8b](#)[70b](#)

↓ 298.1K Pulls ↗ 35 Tags ⌚ Updated 1 year ago

openhermes

OpenHermes 2.5 is a 7B model fine-tuned by Teknium on Mistral with fully open datasets.

↓ 369.7K Pulls ↗ 35 Tags ⌚ Updated 2 years ago

wizard-vicuna-uncensored

Wizard Vicuna Uncensored is a 7B, 13B, and 30B parameter model based on Llama 2 uncensored by Eric Hartford.

[7b](#)[13b](#)[30b](#)

↓ 451.4K Pulls ↗ 49 Tags ⌚ Updated 2 years ago

openchat

A family of open-source models trained on a wide variety of data, surpassing ChatGPT on various benchmarks. Updated to version 3.5-0106.

[7b](#)

↓ 390.8K Pulls ↗ 50 Tags ⌚ Updated 2 years ago

tulu3

Tulu 3 is a leading instruction following model family, offering fully open-source data, code, and recipes by the The Allen Institute for AI.

[8b](#)[70b](#)

↓ 178.9K Pulls ↗ 9 Tags ⌚ Updated 1 year ago

starcoder

StarCoder is a code generation model trained on 80+ programming languages.

1b

3b

7b

15b

↓ 411.3K Pulls ↗ 100 Tags ⌚ Updated 2 years ago

stable-code

Stable Code 3B is a coding model with instruct and code completion variants on par with models such as Code Llama 7B that are 2.5x larger.

3b

↓ 338.8K Pulls ↗ 36 Tags ⌚ Updated 1 year ago

dolphincoder

A 7B and 15B uncensored variant of the Dolphin model family that excels at coding, based on StarCoder2.

7b

15b

↓ 298.7K Pulls ↗ 35 Tags ⌚ Updated 1 year ago

falcon

A large language model built by the Technology Innovation Institute (TII) for use in summarization, text generation, and chat bots.

7b

40b

180b

↓ 395.8K Pulls ↗ 38 Tags ⌚ Updated 2 years ago

nous-hermes

General use models based on Llama and Llama 2 from Nous Research.

7b

13b

↓ 405.1K Pulls ↗ 63 Tags ⌚ Updated 2 years ago

nous-hermes2

The powerful family of models by Nous Research that excels at scientific discussion and coding tasks.

10.7b 34b

↓ 332.7K Pulls ↗ 33 Tags ⌚ Updated 2 years ago

stablelm2

Stable LM 2 is a state-of-the-art 1.6B and 12B parameter language model trained on multilingual data in English, Spanish, German, Italian, French, Portuguese, and Dutch.

1.6b 12b

↓ 314.9K Pulls ↗ 84 Tags ⌚ Updated 1 year ago

vicuna

General use chat model based on Llama and Llama 2 with 2K to 16K context sizes.

7b 13b 33b

↓ 388.7K Pulls ↗ 111 Tags ⌚ Updated 2 years ago

neural-chat

A fine-tuned model based on Mistral with good coverage of domain and language.

7b

↓ 335.7K Pulls ↗ 50 Tags ⌚ Updated 2 years ago

command-r7b

The smallest model in Cohere's R series delivers top-tier speed, efficiency, and quality to build powerful AI applications on commodity GPUs and edge devices.

tools 7b

↓ 133.7K Pulls ↗ 5 Tags ⌚ Updated 1 year ago

smallthinker

A new small reasoning model fine-tuned from the Qwen 2.5 3B Instruct model.

3b

↓ 138.2K Pulls ↗ 5 Tags ⌚ Updated 1 year ago

tinydolphin

An experimental 1.1B parameter model trained on the new Dolphin 2.8 dataset by Eric Hartford and based on TinyLlama.

1.1b

↓ 294.3K Pulls ↗ 18 Tags ⌚ Updated 1 year ago

bge-large

Embedding model from BAAI mapping texts to vectors.

embedding

335m

↓ 186.6K Pulls ↗ 3 Tags ⌚ Updated 1 year ago

llama2-chinese

Llama 2 based model fine tuned to improve Chinese dialogue ability.

7b

13b

↓ 346.2K Pulls ↗ 35 Tags ⌚ Updated 2 years ago

sqlcoder

SQLCoder is a code completion model fined-tuned on StarCoder for SQL generation tasks

7b

15b

↓ 325.9K Pulls ↗ 48 Tags ⌚ Updated 1 year ago

solar-pro

Solar Pro Preview: an advanced large language model (LLM) with 22 billion parameters designed to fit into a single GPU

22b

↓ 166.6K Pulls ↗ 18 Tags ⌚ Updated 1 year ago

solar

A compact, yet powerful 10.7B large language model designed for single-turn conversation.

10.7b

↓ 280.5K Pulls ↗ 32 Tags ⌚ Updated 2 years ago

wizardcoder

State-of-the-art code generation model

33b

↓ 327.6K Pulls ↗ 67 Tags ⌚ Updated 2 years ago

starling-lm

Starling is a large language model trained by reinforcement learning from AI feedback focused on improving chatbot helpfulness.

7b

↓ 282.3K Pulls ↗ 36 Tags ⌚ Updated 1 year ago

mistral-openorca

Mistral OpenOrca is a 7 billion parameter model, fine-tuned on top of the Mistral 7B model using the OpenOrca dataset.

7b

↓ 301.1K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

bespoke-minicheck

A state-of-the-art fact-checking model developed by Bespoke Labs.

7b

↓ 156.8K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

sailor2

Sailor2 are multilingual language models made for South-East Asia.

Available in 1B, 8B, and 20B parameter sizes.

1b

8b

20b

↓ 126.3K Pulls ↗ 13 Tags ⌚ Updated 1 year ago

samantha-mistral

A companion assistant trained in philosophy, psychology, and personal relationships. Based on Mistral.

7b

↓ 294.9K Pulls ↗ 49 Tags ⌚ Updated 2 years ago

xwinlm

Conversational model based on Llama 2 that performs competitively on various benchmarks.

7b

13b

↓ 277.8K Pulls ↗ 80 Tags ⌚ Updated 2 years ago

yarn-llama2

An extension of Llama 2 that supports a context of up to 128k tokens.

7b

13b

↓ 271.1K Pulls ↗ 67 Tags ⌚ Updated 2 years ago

mathstral

Math Σ tral: a 7B model designed for math reasoning and scientific discovery by Mistral AI.

7b

↓ 170.9K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

wizard-math

Model focused on math and logic problems

7b

13b

70b

↓ 299.5K Pulls ↗ 64 Tags ⌚ Updated 2 years ago

orca2

Orca 2 is built by Microsoft research, and are a fine-tuned version of Meta's Llama 2 models. The model is designed to excel particularly in reasoning.

7b

13b

↓ 257.9K Pulls ↗ 33 Tags ⌚ Updated 2 years ago

llama-pro

An expansion of Llama 2 that specializes in integrating both general language understanding and domain-specific knowledge, particularly in programming and mathematics.

↓ 237.8K Pulls ↗ 33 Tags ⌚ Updated 2 years ago

phind-codellama

Code generation model based on Code Llama.

34b

↓ 280.9K Pulls ↗ 49 Tags ⌚ Updated 2 years ago

nuextract

A 3.8B model fine-tuned on a private high-quality synthetic dataset for information extraction, based on Phi-3.

3.8b

↓ 157.4K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

llava-phi3

A new small LLaVA model fine-tuned from Phi 3 Mini.

vision

3.8b

↓ 183K Pulls ↗ 4 Tags ⌚ Updated 1 year ago

meditron

Open-source medical large language model adapted from Llama 2 to the medical domain.

7b

70b

↓ 234.1K Pulls ↗ 22 Tags ⌚ Updated 2 years ago

firefunction-v2

An open weights function calling model based on Llama 3, competitive with GPT-4o function calling capabilities.

tools

70b

↓ 149.2K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

yarn-mistral

An extension of Mistral to support context windows of 64K or 128K.

7b

↓ 234.9K Pulls ↗ 33 Tags ⌚ Updated 2 years ago

granite3-guardian

The IBM Granite Guardian 3.0 2B and 8B models are designed to detect risks in prompts and/or responses.

[2b](#)[8b](#)

↓ 115.3K Pulls ↗ 10 Tags ⌚ Updated 1 year ago

deepseek-v2.5

An upgraded version of DeepSeek-V2 that integrates the general and coding abilities of both DeepSeek-V2-Chat and DeepSeek-Coder-V2-Instruct.

[236b](#)

↓ 131.2K Pulls ↗ 7 Tags ⌚ Updated 1 year ago

stable-beluga

Llama 2 based model fine tuned on an Orca-style dataset. Originally called Free Willy.

[7b](#)[13b](#)[70b](#)

↓ 250.9K Pulls ↗ 49 Tags ⌚ Updated 2 years ago

falcon2

Falcon2 is an 11B parameters causal decoder-only model built by TII and trained over 5T tokens.

[11b](#)

↓ 163.1K Pulls ↗ 17 Tags ⌚ Updated 1 year ago

dbrx

DBRX is an open, general-purpose LLM created by Databricks.

[132b](#)

↓ 170.4K Pulls ↗ 7 Tags ⌚ Updated 1 year ago

nexusraven

Nexus Raven is a 13B instruction tuned model for function calling tasks.

13b

↓ 232.3K Pulls ↗ 32 Tags ⌚ Updated 2 years ago

marco-o1

An open large reasoning model for real-world solutions by the Alibaba International Digital Commerce Group (AIDC-AI).

7b

↓ 103K Pulls ↗ 5 Tags ⌚ Updated 1 year ago

nous-hermes2-mixtral

The Nous Hermes 2 model from Nous Research, now trained over Mixtral.

8x7b

↓ 191.8K Pulls ↗ 18 Tags ⌚ Updated 1 year ago

wizardlm

General use model based on Llama 2.

↓ 230K Pulls ↗ 73 Tags ⌚ Updated 2 years ago

wizardlm-uncensored

Uncensored version of Wizard LM model

13b

↓ 211.7K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

duckdb-nsql

7B parameter text-to-SQL model made by MotherDuck and Numbers Station.

7b

↓ 159.9K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

megadolphin

MegaDolphin-2.2-120b is a transformation of Dolphin-2.2-70b created by interleaving the model with itself.

120b

↓ 157.4K Pulls ↗ 19 Tags ⌚ Updated 2 years ago

medllama2

Fine-tuned Llama 2 model to answer medical questions based on an open source medical dataset.

7b

↓ 191.6K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

stablelm-zephyr

A lightweight chat model allowing accurate, and responsive output without requiring high-end hardware.

3b

↓ 163.9K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

magocoder

 Magocoder is a family of 7B parameter models trained on 75K synthetic instruction data using OSS-Instruct, a novel approach to enlightening LLMs with open-source code snippets.

7b

↓ 162.6K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

notux

A top-performing mixture of experts model, fine-tuned with high-quality data.

8x7b

↓ 151.7K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

notus

A 7B chat model fine-tuned with high-quality data and based on Zephyr.

7b

↓ 151K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

codeup

Great code generation model based on Llama2.

13b

↓ 176.3K Pulls ↗ 19 Tags ⌚ Updated 2 years ago

mistrallite

MistralLite is a fine-tuned model based on Mistral with enhanced capabilities of processing long contexts.

7b

↓ 157.5K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

everythinglm

Uncensored Llama2 based model with support for a 16K context window.

13b

↓ 172.6K Pulls ↗ 18 Tags ⌚ Updated 2 years ago

codebooga

A high-performing code instruct model created by merging two existing code models.

34b

↓ 152.1K Pulls ↗ 16 Tags ⌚ Updated 2 years ago

goliath

A language model created by combining two fine-tuned Llama 2 70B models into one.

↓ 139.1K Pulls ↗ 16 Tags ⌚ Updated 2 years ago

wizard-vicuna

Wizard Vicuna is a 13B parameter model based on Llama 2 trained by MelodysDreamj.

13b

↓ 153K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

open-orca-platypus2

Merge of the Open Orca OpenChat model and the Garage-bAInd Platypus 2 model. Designed for chat and code generation.

13b

↓ 146.5K Pulls ↗ 17 Tags ⌚ Updated 2 years ago

alfred

A robust conversational model designed to be used for both chat and instruct use cases.

40b

↓ 88.5K Pulls ↗ 7 Tags ⌚ Updated 2 years ago
